

# UNIX

---

## Sockets

# Introduction to Sockets

---

- Interprocess Communication channel:
  - descriptor based
  - two way communication
  - can connect processes on different machines
- Three most typical socket types (colloquial names):
  - **unix** (local connection)
  - **tcp** (reliable, byte stream connection over network)
  - **udp** (unreliable, datagram connection over network)

# Introduction to Sockets

---

- Communication styles
  - *SOCK\_STREAM*
    - binary stream
    - reliable
    - context
    - communication with one peer only
    - no records boundaries
    - connection oriented
    - fifo-like connection type
    - bidirectional
    - out of the band data may be (and usually is) supported

# Introduction to Sockets

---

- Communication styles
  - *SOCK\_DGRAM*
    - packets
    - **data loss is possible**
    - **data duplicates are possible**
    - **order of delivery may be different than the order of sending**
    - connectionless
    - communication with many peers
    - each packet has to be individually addressed

# Introduction to Sockets

---

- Communication styles (cont.)
  - *SOCK\_RAW*
    - low level network access
    - e.g. icmp ping request
    - requires proper capability or root privileges
    - SOCK\_DGRAM similar type of connection

# Introduction to Sockets

---

- Addresses, namespace, domain
  - usually **internet** *AF\_INET* or **local (unix)** *AF\_LOCAL*, *AF\_UNIX*
  - new socket can can be bound to an address
    - obligatory for server address
    - client socket is usually bound automatically on first send or connect operation

# Introduction to Sockets

---

- Protocol
  - **tcp** for byte stream contextual reliable connection
  - **udp** for packet contextless unreliable connection
  - usually programs use default (zero) protocol for given communication style and name space
- Connection requires both ends to use the same protocol in the same namespace and communication style

# Introduction to Sockets

---

- Structure `sockaddr` is used for all namespaces:

```
sa_family_t sa_family  
char sa_data[14]
```

- this type is used only to cast addresses used as parameters to `bind`, `getsockname` and other functions
- `sa_family` can be one of *AF\_INET*, *AF\_LOCAL*, *AF\_UNIX*
- each namespace defines its own structure, field `sa_family` is common to all those structures and can be used to recognize socket type



# Introduction to Sockets

---

- If other end of stream socket connection is closed and the connection buffer is empty:
  - reading from this channel will return end of file status (read returns zero as number of bytes read)
  - writing to the channel results in delivery of SIGPIPE signal. If this signal is handled, ignored or blocked `write/send/sendto` will return the EPIPE error
  - writing and reading may result in ECONNRESET error
- EPIPE and ECONNRESET are not critical errors and can be properly recognized and handled in students' apps without terminating the application

# Introduction to Sockets

---

- Socket connection typical use case (very rarely modified)
  - server (connection listener)
    - create socket
    - bind it with the address
    - set listen queue (SOCK\_STREAM only)
    - accept connections (SOCK\_STREAM only)
    - do work
    - close socket
  - client (connection initiator)
    - create socket
    - optionally bind with an address
    - connect
    - do work
    - close

# Introduction to Sockets

---

- Out of Band data
    - priority messages can be send over the socket
    - usually used to indicate exceptional conditions
    - data is sent independently of the ordinary data
    - process must use `send/recv` functions with `MSG_OOB` flag, without this flag only ordinary data will be received
    - `select` and `poll` function can wait for OOB data
    - special marker is set to indicate the position of band data in the ordinary data
    - more information: *glibc manual chapter “Out-of-Band Data”*
-

# Inter-Architectural Socket Connection

---

- Sockets can be used to communicate processes working on different architectures, this affects the way data types are treated
- **Failure to recognize the problem is the most common source of errors in students' works**
- Byte Order
  - does not affect 1 byte data types
    - `char` (not unicode `wchar_t`)
    - `[u]int8_t`
    - the safest method to send data is to format it as text !

# Inter-Architectural Socket Connection

---

- Byte Order (cont.)
  - 2- and more- byte integers
    - little-endian less significant byte goes first
      - the higher memory address the more significant byte (255=FF00)
      - x86, amd64
    - big-endian more significant byte goes first
      - the lower memory address the more significant byte (255=00FF)
      - SPARC, PPC
    - network byte order (big-endian), all data should be converted to network byte order before being sent and back to host byte order after being received

# Inter-Architectural Socket Connection

---

- Byte Order (cont.)
  - 2 byte data type (`[u]int16_t`) can be converted to network order with `htons` function and back to host order with `ntohs` function
  - 4 byte data type (`[u]int32_t`) can be converted to network order with `htonl` function and back to host order with `ntohl` function
  - 8 byte data types (`[u]int64_t`) must be converted manually, no standard function or macro exists, it is easy to write one though

# Inter-Architectural Socket Connection

---

- Byte Order (cont.)
  - double and float
    - do not have network format defined !
    - the floating point format may be different !
    - it should be sent as human readable string and parsed at destination
  - arrays
    - all members of array should be converted separately
  - structures
    - all members of a structure should be converted separately
    - structure must be packed (see next slides)

# Inter-Architectural Socket Connection

---

- Data types sizes
  - `char` is a one-byte data type (always)
  - do not use traditional data types `int` and `long` as they may differ in size
  - use UNIX standardized integer types:
    - `[u]int[8|16|32|64]_t`
- Structures have different field alignment
  - some architectures prefer 2,4 or 8 byte alignment
  - if communicating architectures use different alignment the structure will be malformed after transport



# Inter-Architectural Socket Connection

---

- Structures have different field alignment (cont.)
  - compiler can enforce the smallest alignment possible that is the same on all architectures – so called *packing*
  - unfortunately packing options are specific to given compilers
    - some use `#pragma pack;`
    - gcc uses `__attribute__((__packed__))` on all structures to pack or `-fpack-struct` option to pack all structures in the code
  - **there is no portable way of packing, structures should not be sent over socket in portable programs**

# Socket Functions

---

- Function `socket` :
  - creates unbound and unconnected socket descriptor
  - `domain` selects namespace, usually one of *PF\_INET*, *PF\_LOCAL*, *PF\_UNIX* (**notice that constant has *PF\_* not *AF\_* prefix**)
  - `type` indicates communication style, usually *SOCK\_STREAM* or *SOCK\_DGRAM*
  - `protocol` usually defaults to zero as only one protocol exists for each combination of `domain` and `type`

# Socket Functions

---

- Function `socketpair` :
  - creates a pair of unnamed and connected socket
  - only related processes can use it
  - can be used the same way as pipe, except that it offers two-way communication
  - domain can be only *PF\_LOCAL*, *PF\_UNIX* on Linux and most of other systems
  - other parameters are passed similarly as to socket function
  - `socket_vector` – array of 2 undistinguishable socket descriptors (return parameter)

# Socket Functions

---

- Function `bind` :
  - assigns name to the socket
    - not obligatory, process can use unbound socket, and the name will be assigned automatically, process can use `f.getsockname` to learn this name
    - name of local end of the connection is assigned
    - name of remote peer can be read with `f.getpeername`
  - address must be cast to `struct sockaddr *`
  - address structure real length must be calculated accordingly to namespace type (see examples) as it is not always matter of a simple `sizeof`.

# Socket Functions

---

- Functions `setsockopt` and `getsockopt`:
  - control various socket options (see examples)
  - useful POSIX compliant options for `AF_INET`:
    - `SO_REUSEADDR`(`SOL_SOCKET` level)
      - application fails to bind to an address some time after previous run (no matter how it was terminated). It takes even a few minutes for the system to clear the address, this option will speed up reuse, it should be instant but makes tcp less reliable as time-out for lost packets is removed
    - `SO_KEEPALIVE`(`SOL_SOCKET` level)
      - starts transparent exchange of test message, if connection is broken it can speed up detection of the problem. Waiting processes are sent `SIGPIPE` if test message fails to go through.

# Socket Functions

---

- Functions `setsockopt` and `getsockopt(cont.)`:
  - `SO_BROADCAST` (`SOL_SOCKET` level)
    - permits sending broadcasts from the sockets
    - broadcast can be blocked on firewall, usually no broadcasts from the external network are allowed in LAN
  - `IP_MTU` (`IPPROTO_IP` level)
    - reads current known MTU (Maximum Transport Unit) - the maximum reliable size of datagram process can send
    - can be obtained (never set) for connected socket only
    - depends on network connection
    - can change
    - cannot be less than 576
  - `SO_ERROR` (`SOL_SOCKET` level)
    - check for network pending errors

# Socket Functions

---

- Function `connect` :
  - connects socket with the other end
  - usually client connects to the server
  - socket being connected does not need to be bound to any name
  - server socket must have a name

# Socket Functions

---

- Function `connect` (cont.):
  - packet (datagram)
    - does not require permanent connection as each packet can be addressed individually
    - connecting is still possible and useful when many packets are to be sent to the same destination
    - process can connect the same packet socket many times to change destination address
    - `connect` will never block as the real connection is not established



# Socket Functions

---

- Function `connect` (cont.):
  - `SOCK_STREAM` connection
    - connection is obligatory and possible only once per given socket
    - `connect` will block (provided that `O_NONBLOCK` flag was not set on socket descriptor) until connection is made
    - if socket is in non-blocking mode, or `connect` was interrupted by signal handling routine, connecting will **continue asynchronously !!!** In such a case:
      - process should wait for socket write readiness (see examples)
      - or assume socket to be connected, sleep 1-2 sec. If assumption is wrong the first operation on socket will fail (rather untidy approach)
      - repeated connect returns `EALREADY` error, it must be recognised if `TEMP_FAILURE_RETRY` macro is used

# Socket Functions

---

- Function `listen` :
  - applies only to byte stream communication
  - enables connection requests on the socket
  - marks socket as server socket that should be used only to accept connections
  - `backlog` argument
    - merely a hint of size of queue of pending connections (i.e. connections already initiated by client and not yet accepted)
    - **not** the maximum number of simultaneously connected clients
    - POSIX suggests that passing zero as backlog parameter may mean implementation-defined minimal value
    - if pending connections queue is exhausted peer may be rejected

# Socket Functions

---

- Function `accept` :
  - applies only to byte stream communication
  - creates a new socket for communication with connecting peer, returns new descriptor
  - new socket communication type and namespace are inherited from listening socket
  - newly created socket may not inherit flags set on listening socket (on Linux flags are not inherited), portable application should set flags explicitly
  - listening socket remains unconnected and can accept next connection
  - connections are accepted in the order they were queued

# Socket Functions

---

- Function `accept` (cont.):
  - blocks unless non-blocking flag is set for descriptor
  - listening socket indicates readiness for a new connection by readable event (f. `select`), even when `select` returns, such a connection may be already lost due to network errors, thus `accept` will block. To avoid this situation, non-blocking mode should be set on the socket.

# Socket Functions

---

- Function `accept` and `getpeername` (cont.):
  - can obtain peer socket address
    - via output parameters
    - if address buffer is too small, the address will be **silently** truncated (the function cannot overflow the buffer)
    - if connection is made by unbound peer the value of the address is unspecified (udp and tcp use bounded clients only)
    - there is a very interesting discussion on the type of last argument to `accept` (see NOTES in Linux manual on the `accept` function)
    - On HP-UX does not support `AF_UNIX`

# Socket Functions

---

- Functions `socket`, `socketpair`, `listen`, `accept` and `bind` (cont.):

- from linux manual:

*POSIX.1-2001 does not require the inclusion of `<sys/types.h>`, and this header file is not required on Linux. However, some historical (BSD) implementations required this header file, and portable applications are probably wise to include it.*

# Socket Functions

---

- Operations on socket
  - blocks if operation buffer is full (and flag `O_NONBLOCK` is not set)
  - success of operation does not mean the delivery, only successful sending
  - if datagram is received it must be read at once, otherwise the remaining part will be truncated !
  - can report errors from previous operations (pending network errors)
  - using `read/write` functions on connected socket

# Socket Functions

---

- Operations on socket (cont.)
  - with `recv/send` functions on connected socket
    - works as `read/write`
    - process can apply socket specific flags (POSIX list)
      - `MSG_OOB` – out of band data
      - `MSG_EOR` – end of record marker (not supported on byte stream)
      - `MSG_PEEK` – do not remove data from the buffer
      - `MSG_WAITALL` – the read shall not be terminated in the middle of a message (but still can be interrupted before any data is actually read)
      - flag `MSG_DONTWAIT` is Linux specific !
      - flag `MSG_NOSIGNAL` is Linux specific !



# Socket Functions

---

- Operations on socket (cont.)
  - with `recvfrom/sendto` functions
    - works as `recv/send`
    - receiving process can learn peer address
    - sending process can address the message (datagrams only)
  - with `recvmsg/sendmsg` functions
    - works as `recvfrom/sendto`
    - sends messages with control information

# Socket Functions

---

- Closing a socket
  - with `close` function
  - with `shutdown` function a process can shut only a part of a connection
    - `SHUT_RD` – process stops reading from the socket any incoming data should be rejected, process can still send data. If process tries to read from such a socket it will get EOF marker. Peer may get EPIPE or data may be silently discarded, in practice data is still received and queued, can be read in some systems !
    - `SHUT_WR` - process stops writing to socket any outgoing data will be rejected, but process still can read from the socket. If process tries to write to such a socket it will be send SIGPIPE (or EPIPE error)
    - `SHUT_RDWR` – closes both ways, still requires `f.close` .

# Unix Domain Sockets

---

- Communication between processes on the same computer
  - socket name is simply a name in local file system
  - processes must have write (w) and search (x) permissions to the directory where socket is created
  - process must have read and write permissions to the directory containing the socket it is connecting to
  - permissions on local sockets can be ignored by some systems and should not be used as security measures
  - usually local sockets are put in `/tmp`

# Unix Domain Sockets

---

- both communication styles (*SOCK\_STREAM*, *SOCK\_DGRAM*) are supported in this namespace
- name in file system is necessary only to establish connection
- socket file persists and can/should be deleted by the process
- local sockets can not be used to connect remote process (even through nfs or afs) at the moment but it is possible that it will change in the future.

# Unix Domain Sockets

---

- for better portability process should not depend on the fact that the other process resides on the same machine
  - pid numbers
  - byte order
  - sending structures
- for sake of simplicity students can assume that local sockets are truly local in laboratory applications
- Linux implementation does not support out of bound data transfer on local sockets

# Unix Domain Sockets

---

- Address in unix domain sockets are stored in `struct sockaddr_un`:  
  
`sa_family_t sun_family`  
`char sun_path[108]`
- The length of this structure (passed to `bind`) should be calculated as:

`sizeof(sa_family_t) + strlen(a.sun_path)`

- The `SUN_LEN` macro returns proper size
- **Do not use `sizeof(sockaddr_un)`!!!**

# Internet Domain Sockets

---

- Communication between processes over the tcp/ip network
  - IPv4 vs. IPv6, at the moment only IPv4 is common and will be discussed here
  - socket name consists of IP address (x.x.x.x) and the port number **both represented in network byte order** (see byte order)
  - both communication styles (*SOCK\_STREAM*, *SOCK\_DGRAM*) are supported in this namespace
  - **no permission control built into AF\_INET sockets !!!**

# Internet Domain Sockets

---

- Address in internet domain sockets are stored in `struct sockaddr_in`:

```
sa_family_t sin_family
struct in_addr sin_addr
unsigned short int sin_port
```

- Where `in_addr` is a one-filed structure:

```
u_int32_t s_addr;
```

- The length of this structure (passed to `bind`) should be calculated as

```
sizeof(sockaddr_in)
```



# Internet Domain Sockets

---

- Special addresses (host byte order):
  - `uint32_t INADDR_LOOPBACK (127.0.0.1)`
  - `uint32_t INADDR_ANY (0.0.0.0)`
    - very useful as server address
    - binds to all local interfaces
    - if socket is unbound prior to connect or send it will be automatically bound to `INADDR_ANY` and random free port
  - `uint32_t INADDR_BROADCAST (255.255.255.255)`
    - to send broadcast datagrams (`SOCK_DGRAM`)
  - `uint32_t INADDR_NONE (-1==255.255.255.255)`
    - error indicator

# Internet Domain Sockets

---

- IP address by type:
  - unicast – addressing of single peer
    - only option for SOCK\_STREAM
  - multicast – addressing a set of peers in the network
    - not available by default, popular, speeds up streaming
    - will not be discussed here
  - broadcast – addressing all the peers in the network
    - useful communication method in local networks
    - can be blocked by routers and local firewalls and we experienced many problems with broadcast tasks during labs

# Internet Domain Sockets

---

- Socket address (cont.)
  - to convert IP string e.g. “194.29.178.1” to `uint32_t`  
POSIX defines `inet_addr` function
    - returns `INADDR_NONE` as error, it may be mistaken with `INADDR_BROADCAST`
    - `inet_addr` returns an address in **network byte order**
    - GNU extension `inet_aton` reports errors more reliably but is not standardized
  - to convert binary address in **network byte order** to string  
POSIX defines `inet_ntoa`
    - returns string in static buffer – not thread safe

# Internet Domain Sockets

---

- Socket address (cont.)
  - as socket requires IP number to bind, domain name must be resolved (DNS) with `gethostbyname`
    - on success returns `hostent` structure, where address in question is available as field `char *h_addr_list[0]` (must be cast to `struct in_addr`)
    - errors are reported in global variable **`h_errno`** not `errno` !!!
    - can be used with domain names and IP numbers as strings on Linux (POSIX does not define it)
    - returns addresses in network byte order
    - returned address may be stored in static buffer
    - not thread safe

# Internet Domain Sockets

---

- Datagrams (packets)
  - if size of datagram is larger than MTU on given socket sending will return with EMSGSIZE error
  - process can test MTU with `getsockopt` or experiment with the sizes and EMSGSIZE or assume small datagrams (less than 576)
  - `read/write/send/rcv/sendto` and `recvfrom` will perform atomically on datagrams
  - if datagram is received it must be read at once, otherwise the remaining part will be truncated !

# Internet Domain Sockets

---

- Datagrams (packets) cont.
  - if datagram is lost, application must retransmit it, usually if response is not received after some time-out packet can be considered lost
  - usually messages are sent in single packets
  - application logic must be ready for lost packets and duplications
  - on Linux datagrams are reliable, duplications and mixed order or delivery are not possible, but this is not POSIX behaviour
  - Linux also has a reliable datagram SOCK\_RDM socket type

# Internet Domain Sockets

---

- Binary stream
  - reliable
  - usually only one process operates on each end of connection
  - message boundaries are not preserved
  - **there is no atomic message size !!!**
  - **read/write operation can be interrupted at any stage (EINTR)**

# Internet Domain Sockets

---

- Binary stream (cont.)
  - usually connection is bidirectional (read/write), process can modify file flags or use `shutdown` function to limit access
  - file position is fixed, reading from beginning, writing at the end (FIFO order)
  - read buffer is independent from write buffer



# Sockets - Examples

---

- How to make local SOCK\_STREAM socket :

```
int make_socket(char* name , struct sockaddr_un \
*addr) {
    int sockfd;
    if ((sockfd = socket(PF_UNIX, SOCK_STREAM, \
                        0)) < 0) ERR("socket");
    memset(addr, 0, sizeof(struct sockaddr_un));
    addr->sun_family = AF_UNIX;
    strncpy(addr->sun_path, name, sizeof( \
                                           addr->sun_path)-1);
    return sockfd;
}
```

# Sockets - Examples

---

- How to bind and start listening on local SOCK\_STREAM socket :

```
#define BACKLOG 3
```

```
int bind_socket(char *name) {  
    struct sockaddr_un addr;  
    int sockfd;  
    if (unlink(name) < 0 && errno != ENOENT)  
        ERR("unlink");  
    sockfd = make_socket(name, &addr);  
    if (bind(sockfd, (struct sockaddr*) &addr, \  
        SUN_LEN(&addr)) < 0) ERR("bind");  
    if (listen(sockfd, BACKLOG) < 0) ERR("listen");  
    return sockfd;  
}
```

# Sockets - Examples

---

- How to connect to local SOCK\_STREAM socket :

```
int connect_socket(char *name) {
    struct sockaddr_un addr; int sockfd;
    sockfd = make_socket(name, &addr);
    if (connect(sockfd, (struct sockaddr*) &addr, \
                SUN_LEN(&addr)) < 0) {
        if (errno != EINTR) ERR("connect");
        else {
            fd_set wfds ; int status;
            socklen_t size = sizeof(int);
            FD_ZERO(&wfds); FD_SET(sockfd, &wfds);
            if (TEMP_FAILURE_RETRY(select(sockfd+1, \
                NULL, &wfds, NULL, NULL)) < 0) ERR("select");
            if (getsockopt(sockfd, SOL_SOCKET, SO_ERROR, \
                &status, &size) < 0) ERR("getsockopt");
            if (0 != status) ERR("connect");
        }
    }
    return sockfd; }
```

---

# Sockets - Examples

---

- How to accept connection on SOCK\_STREAM socket :

```
int add_new_client(int sfd, fd_set* base_rfds, \
                  int* fdmax) {
    int fd;
    if ((fd=TEMP_FAILURE_RETRY(accept(sfd, NULL, \
        NULL))) < 0) ERR("accept");
    FD_SET(fd, base_rfds);
    *fdmax = (*fdmax < fd ? fd : *fdmax);
    return 1;
}
```

# Sockets - Examples

---

- How to make internet server SOCK\_DGRAM socket :

```
int make_socket(uint16_t port) {  
    struct sockaddr_in name;  
    int sock, t=1;  
    sock = socket(PF_INET, SOCK_DGRAM, 0);  
    if (sock < 0) ERR("socket");  
    name.sin_family = AF_INET;  
    name.sin_port = htons(port);  
    name.sin_addr.s_addr = htonl(INADDR_ANY);  
    if (setsockopt(sock, SOL_SOCKET, SO_REUSEADDR,  
        &t, sizeof(t))) ERR("setsockopt");  
    if (bind(sock, (struct sockaddr*) &name, \  
        sizeof(name)) < 0) ERR("bind");  
    return sock;  
}
```

---

}

# Sockets - Examples

---

- How to receive datagrams:

```
int16_t recv_datagram(int sock, struct sockaddr_in
*addr) {
    int16_t buf;
    socklen_t len = sizeof(struct sockaddr_in);
    if (TEMP_FAILURE_RETRY(recvfrom(sock, &buf, \
        sizeof(int16_t), 0, (struct sockaddr*) addr, \
&len)) \ < sizeof(int16_t))
        ERR("recvfrom");
    return ntohs(buf);
}
```

# Sockets - Examples

---

- How to send datagrams:

```
int send_datagram(int sock, struct sockaddr_in addr,
int16_t msg) {
    int status;
    int16_t buf = htons(msg);
    status = TEMP_FAILURE_RETRY(sendto(sock, &buf, \
        sizeof(int16_t), 0, (struct sockaddr*) \
        &addr, sizeof(addr)));
    if (status < 0 && errno != EPIPE && errno ==
ECONNRESET)
        ERR("sendto");
    return status;
}
```

# Sockets - Examples

---

- How to lookup domain name in DNS:

```
struct sockaddr_in make_address(char *address,
uint16_t port) {
    struct sockaddr_in addr;
    struct hostent *hostinfo;
    addr.sin_family = AF_INET;
    addr.sin_port = htons (port);
    hostinfo = gethostbyname (address);
    if (hostinfo == NULL) ERRH("gethost:");
    /*h_errno*/
    addr.sin_addr =
        * (struct in_addr *) hostinfo->h_addr;
    return addr;
}
```